

The ad evidence layer

It reads a company's own website and answers one question honestly: *are they buying traffic?* It never answers *how much* — because no page on any website states a media budget, and a number we invented would look exactly like a number we measured.

WHAT IT DOES

It turns markup into a qualification verdict

When we crawl a prospect's site for ICP classification, we already have the homepage HTML in hand. This layer makes one extra pass over that markup — no extra fetch, no extra API call, no per-lead cost — looking for the tags a company installs when it runs paid campaigns.

The distinction that carries the whole thing is between a **measurement** tag and a **conversion** tag. Google Analytics means a company looks at its traffic. A Google Ads conversion ID means a company is paying for traffic and wants the platform to optimise against a purchase. Nobody installs the second one speculatively.

Everything is deterministic. Same HTML in, same verdict out. The thresholds are constants in a file, not sentences in a prompt — so changing where the line sits is a code review, not a vibe.

01

Crawl

The homepage fetch we already do for ICP scoring. Nothing new is requested.

02

Read

24 tag fingerprints, conversion IDs, pixel IDs, and paid landing-page paths.

03

Score

Tag weight, plus platform breadth, plus campaign infrastructure.

04

Classify

One of four states, with a confidence figure about the classification.

05

Confirm by hand

Deep links straight into the Meta and Google ad libraries for the shortlist.

WHY STEP 05 STAYS HUMAN

Lucas said it himself: *"just having a view/count of the ads that are actively running doesn't really paint the full picture."* Neither library offers an API that gives US ad volume, so the count is a human action either way. What we can do is remove the searching. The verdict ships with the exact library URLs for that brand, so a two-minute lookup becomes a five-second click.

THE VERDICT

Four states, and the fourth is the one that matters

Three states describe what we found. The fourth describes what we never looked for — and keeping it separate from "found nothing" is the difference between a working funnel and one that quietly deletes good prospects.

CONFIRMED

A conversion tag is installed and the surrounding evidence backs it up. They are buying traffic and optimising against outcomes.

score $\geq$ 5 AND a conversion tag

PROBABLE

Pixels are present but no conversion tag. Real enough for a shortlist, not enough for a claim — a pixel can outlive its last campaign by years.

score $\geq$ 2

NO EVIDENCE

We looked at the markup and found no ad platform tags, no pixels, no campaign landing pages.

score < 2

NOT ASSESSED

The crawl never looked — the in-house crawler served this lead, and it doesn't read ad tags. Not a clean bill of health.

no signals collected

THE FAILURE THIS PREVENTS

Collapse "not assessed" into "no evidence" and every lead the fallback crawler touched looks like a confirmed non-advertiser. You would filter out good prospects and never see it happen, because the funnel would report a clean, plausible result.

WHAT THE CRAWLER SEARCHES FOR

Twenty-four fingerprints, weighted by what they prove

Each pattern is matched against every script, link and iframe source on the page, and against inline script text — so a self-hosted embed is still caught. Weight is not popularity; it is how hard the tag is to explain away.

FINGERPRINT IN THE MARKUP	PLATFORM	WT	WHAT IT PROVES
WEIGHT 3 — CONVERSION TRACKING · INSTALLED FOR ONE REASON: TO OPTIMISE PAID SPEND			
googleadservices.com/pagead/conversion	Google Ads	3	Conversion tracking
gtag('config', 'AW-...')	Google Ads	3	Conversion tracking
/pagead/viewthroughconversion	Google Ads	3	View-through conversion
google_conversion_id	Google Ads	3	Conversion tracking
fbq('track', 'Purchase')	Meta	3	Purchase conversion event
WEIGHT 2 — PIXELS AND AUDIENCE TAGS · STRONGLY SUGGEST CAMPAIGNS, CAN BE STALE			
googleads.g.doubleclick.net	Google Ads	2	Remarketing tag
connect.facebook.net/.../fbevents.js	Meta	2	Meta Pixel
fbq('init', ...)	Meta	2	Meta Pixel
analytics.tiktok.com · ttq.load	TikTok	2	TikTok Pixel

FINGERPRINT IN THE MARKUP	PLATFORM	WT	WHAT IT PROVES
bat.bing.com/bat.js · uetq	Microsoft Ads	2	UET tag
snap.lidn.com/li.lms-analytics	LinkedIn Ads	2	Insight Tag
s.pining.com/ct/core.js · pintrk(	Pinterest Ads	2	Pinterest Tag
sc-static.net/scevent	Snapchat Ads	2	Snap Pixel
static.criteo.net	Criteo	2	Retargeting tag
WEIGHT 1 — SUPPORTING STACK · MEASUREMENT MATURITY, NOT PROOF OF SPEND			
triplewhale · triplepixel	Triple Whale	1	Paid attribution platform
northbeam	Northbeam	1	Paid attribution platform
hyros	Hyros	1	Paid attribution platform
static.klaviyo.com · klaviyo.js	Klaviyo	1	Lifecycle marketing
cdn.attn.tv	Attentive	1	SMS marketing
gtm.js?id=GTM-...>m_auth	Server-side GTM	1	Advanced measurement

Two more things it pulls out

Identifiers. The literal `AW-1234567` conversion ID and the Meta pixel ID. These are worth keeping because they are stable, checkable handles — the same ID across two domains tells you it's one operator running both.

Campaign landing pages. Paths that only exist to receive bought traffic: `/lp/`, `/offer`, `/promo`, `/campaign`, `/try-`, `/free-trial`. A brochure site has `/about` and `/contact`. A site running campaigns has `/lp/spring-sale`. The count is an infrastructure signal, not a budget.

How the score is built

The fingerprint weights are summed, then two adjustments: **+1 per additional platform** beyond the first, capped at +3 — running three platforms is a managed programme, not a boosted post; and **+1 for any campaign landing page, +2 for three or more**. Every point is itemised in the output, so a threshold argument stays falsifiable.

RELIABILITY

What you can actually lean on

Not all of this evidence is equally good, and pretending otherwise is how a qualification layer loses trust the first time it's wrong in front of a client.

Google Ads conversion ID RELY ON IT

An `AW-` conversion tag exists solely to feed Google's bidding algorithm. It is not installed by accident, not left behind by a template, and not part of any analytics starter kit. If it's on the page, someone is running Google Ads and paying attention to what those ads produce.

Meta Pixel

RELY WITH CARE

Extremely common, and it can sit on a site for years after the last campaign stopped. Agencies also install it for retargeting alone. It earns a shortlist place; it does not earn a verdict. This is exactly why a lone pixel scores *probable*, never *confirmed*.

Attribution software

BEST SCALE PROXY

Triple Whale, Northbeam and Hyros are paid subscriptions bought by companies whose ad spend is large enough to justify attribution software. It is the closest thing to a spend indicator that is actually observable on a website — which is why it is reported under *scale indicators* and never converted into a dollar figure.

GA4 and plain GTM

DON'T LEAN ON IT

Measurement, not spend. Nearly every site above a certain size has both. They contribute nothing to the score and they never should.

HONEST LIMITS

Three ways this under-detects — and they hit the biggest advertisers hardest

All three failure modes are false negatives, and they skew toward sophisticated advertisers. That inverts the usual intuition: a *no evidence* result on a large brand is more likely a detection gap than a fact about the brand.

1. Tags injected by a container

When Google Ads conversion tracking is deployed through a GTM container, the raw HTML shows only `gtm.js` — the `AW-` ID lives inside the container, fetched and executed at runtime. We see it only if the page was rendered in a browser. Today the crawler renders when a page is blocked or too thin, not when its text is already sufficient. So the tag-rich, text-rich sites are precisely the ones we may read pre-render.

2. Server-side tagging

Meta's Conversions API and server-side GTM move the tag off the page entirely. The companies most likely to have made that move are the ones with enough spend to care about signal loss and iOS restrictions — the \$10K+/month segment we're trying to find.

3. Consent gating

Tags held behind a consent manager may not appear until a visitor accepts. Usually the script reference is still in the source, but not always.

THE RULE THAT FOLLOWS FROM THIS

No evidence must never be a hard disqualifier. It's a sorting signal — it pushes a lead down the queue, it does not remove it. The state exists to rank a list, not to reject a company.

WHAT I COULD NOT MEASURE

I wanted to give you a real hit rate across known advertiser domains, but outbound requests from this environment are proxied and every fetch failed. The honest number has to come from a calibration run on a machine with real egress — ideally against Jeremy's twenty known-spend clients, where we already know the right answer and can set the thresholds against it instead of guessing.

WHAT SHIPPED

Changes in the repo

Twenty tests pass — thirteen on the classifier, seven on the API contract. No database migration: the verdict lives in the lead's existing JSON column.

<code>app/enrichment/ad_signals.py</code>	NEW The classifier, the thresholds, the verify links, the export columns.
<code>app/enrichment/company_research.py</code>	Two bug fixes — see below. This is the important one.
<code>app/enrichment/pipeline.py</code>	Stores the verdict and raw signals on every lead it processes.
<code>app/routers/enrich_lists.py</code>	Exposes the verdict to the UI; adds six columns to the CSV export.
<code>frontend/.../EnrichListDetail.jsx</code>	An "Ad activity" panel in the research drawer, with the library links.
<code>tests/test_ad_signals.py</code>	NEW 13 tests pinning the policy, not the plumbing.
<code>tests/test_company_research_contract.py</code>	NEW 7 tests pinning the shape of the API response.

FOUND WHILE WIRING THIS UP

The Company Research API integration had never actually run. The mapper read each page body from `markdown` or `text`; the API has always called that field `content` — `output_format` changes the encoding of the string, never the key. So every page mapped to an empty string, every response looked like a failed crawl, and the pipeline silently fell back to the in-house crawler. Nothing errored. Nothing logged. Every enrichment you have run so far was in-house-crawler quality. Fixed, and the contract test now fails against the old mapper so it can't come back.

NEXT

What would make this materially better

- 01 Force a render on qualification runs.** This is the single biggest accuracy win available. The scrape endpoint already supports `render: always`; the crawl endpoint has no force-render flag yet. Adding one closes failure mode 1.

- 02 **Calibrate against Jeremy's known-spend clients.** Twenty companies where the answer is already known turns every threshold in this document from a judgement call into a measurement.
- 03 **An "Advertisers" view in the list rail,** so a list can be sorted down to the shortlist in one click rather than by opening leads one at a time.
- 04 **Read Lucas's script when it arrives Monday.** It encodes the questions Unicorn actually trusts for confirming spend in conversation. This layer should feed that script, not duplicate it.

Built for RevCadence · classifier is deterministic and versioned in the repo · no spend figure is produced anywhere in this system